D2: How to Argue Well

A guide to perspective-taking arguments for FIT1061

Your student ID selects a motion and a side. Your job is to build the strongest possible case, even if you disagree with it. This guide shows you how.

---

1. Steel-manning vs straw-manning

Straw-manning is presenting a weak version of a position so it's easy to knock down. **Steel-manning** is presenting the *strongest* version of a position: the version a genuine advocate would make.

	Straw-man	Steel-man
What it sounds like	"People who think AI should make bail decisions just don't care about fairness."	"Advocates for AI in bail decisions argue that human judges are also biased (influenced by fatigue, mood, and implicit prejudice) and that a well-audited algorithm could at least make that bias visible and measurable."
Why it matters	Easy to refute but unconvincing. The audience knows you're not engaging honestly.	Harder to argue, but far more persuasive. Shows you understand the issue deeply enough to argue both sides.

The test: Could someone who genuinely holds this position watch your video and say, "Yes, that's exactly what I believe and why"? If yes, you've steel-manned. If they'd say "That's a caricature of my view," you've straw-manned.

---

2. Using evidence specifically

The difference between a vague argument and a strong one is *specificity*. Compare:

Vague: "AI systems can be unfair to certain groups."

Specific: "In 2016, ProPublica found that COMPAS (a tool used in US courts to predict the tendency of convicted individuals to reoffend) had a false positive rate of 45% for Black defendants compared to 23% for white defendants. Two people with identical risk profiles were treated differently based on race."

The specific version names the system (COMPAS), the context (US courts), the metric (false positive rate), and the numbers (45% vs 23%). It gives the audience something concrete to engage with.

Where to find evidence for your argument: - Your task notebooks (P2–P9, C1–C4): specific results you computed - Milestone case studies from workshops (Deep Blue, AlphaGo, COMPAS, Gender Shades, ImageNet) - Kialo debates (where your motion's week ran one): arguments peers made that you found compelling or challenging - Course readings: Model Cards for Model Reporting (Mitchell et al.) and Green AI (Schwartz et al.)

---

3. Engaging with counterarguments

The weakest part of most student arguments is the counterargument section. Here's a framework:

The acknowledge–pivot–respond pattern

1. **Acknowledge** the counterargument's strength: "The strongest objection to my position is…"
2. **Pivot** by explaining what it gets right: "This is a fair point because…"
3. **Respond** by explaining why your thesis still holds: "However, my position stands because…"

Weak counterargument engagement: > "Some people might say AI isn't that bad. But they're wrong because AI is biased."

Problems: picks a weak objection ("isn't that bad"), dismisses it without engaging ("they're wrong"), and responds with a vague claim.

Strong counterargument engagement: > "The strongest argument against requiring Model Cards is that they create compliance overhead without changing behaviour: companies could publish technically accurate cards that still obscure real limitations. That's a genuine concern. But the alternative (no documentation at all) is worse. Model Cards don't solve the problem, but they make the problem *visible*. And visibility is the precondition for accountability."

This version picks a *strong* objection (compliance theatre), acknowledges it's legitimate, and responds substantively without dismissing it.

---

4. Annotated exemplar argument

Below is a complete argument outline for **Motion 1, arguing AGAINST** ("Algorithmic decision-making is NOT more fair than human decision-making"). Read the annotations in **[brackets]** to see what each section does.

Opening (~30 seconds)

> "The motion asks whether it's acceptable for an AI system to have different error rates for different demographic groups, as long as overall accuracy is high. I'm arguing that it is not, and I want to show you why 'overall accuracy' is the wrong metric when the errors fall unevenly on the people who can least afford them."

[States the motion clearly. Previews the thesis. Signals the core tension: aggregate vs individual impact.]

Point 1 (~60 seconds)

> "In Week 9, we studied Gender Shades, a 2018 study that tested three commercial facial recognition systems. Overall accuracy was above 90% for all three. But when Buolamwini and Gebru broke it down by gender and skin type, the gap was staggering: error rates for light-skinned men were below 1%, while error rates for dark-skinned women reached 34%. The 'overall' number hid a 34-fold disparity."

[Names the study, the researchers, the year, the specific numbers. Grounds the abstract claim in concrete evidence from the course.]

Point 2 (~60 seconds)

> "In C3, we implemented two fairness metrics (demographic parity and equalised odds) and discovered they can contradict each other. You can't satisfy both simultaneously. That's not a bug; it's a mathematical impossibility proven by Chouldechova and others. So when someone says 'overall accuracy is high,' they're choosing one way to measure fairness and ignoring others. The question isn't whether the system is fair: it's *which definition of fairness* was chosen, by whom, and who bears the cost."

[References a specific task (C3) and a specific result (fairness impossibility). Connects the evidence to the thesis: "overall accuracy" is a choice, not a neutral fact.]

Point 3 (~60 seconds)

> "Consider what these error rates mean in practice. If a hiring algorithm rejects qualified candidates from one demographic group at twice the rate of another, those individuals lose job opportunities, not because they're unqualified, but because the training data reflected historical patterns of exclusion. The harm isn't hypothetical; it's the Amazon recruiting tool, shelved in 2018 after it was found to systematically downgrade resumes that included the word 'women's'."

[Moves from metrics to real-world impact. Names a specific case (Amazon hiring tool, 2018). Connects the abstract fairness discussion to concrete human consequences.]

Counterargument + Response (~45 seconds)

> "The strongest argument for accepting different error rates is pragmatic: equalising error rates across groups may require reducing accuracy for some groups, meaning more total errors overall. If the system is making medical diagnoses, more total errors means more missed cancers. That's a serious concern.
>
> But I'd argue it proves my point rather than refuting it. If we have to choose between aggregate accuracy and equitable accuracy, that's a *value judgement*, not a technical decision. And value judgements about who bears the cost of errors should be made transparently (by policymakers, ethicists, and affected communities) not hidden inside an optimisation function."

[Picks the STRONGEST counterargument (total error trade-off). Acknowledges it's legitimate ("That's a serious concern"). Responds by reframing: the trade-off exists, but it proves the decision is ethical, not technical.]

Closing (~15 seconds)

> "Overall accuracy is a single number that hides who is being harmed. Until we can guarantee that errors are distributed fairly (or at least transparently) different error rates for different groups should not be acceptable."

[Restates thesis. Ends cleanly.]

---

5. Good vs weak, side by side

Here are two openings for **Motion 1, arguing AGAINST** ("Algorithmic decision-making is NOT more fair than human decision-making"). Compare:

Weak opening: > "I'm arguing that algorithms aren't fair. I think this is really important because AI can be biased and that's not fair to people."

Problems: - Signals personal opinion ("I think") instead of building a case - Vague ("can be biased"): no specifics - No preview of the argument structure

Strong opening: > "In 2016, a piece of software called COMPAS was used in courtrooms across the United States to help judges assess the tendency of convicted individuals to reoffend. ProPublica's investigation found that Black defendants were nearly twice as likely to be incorrectly flagged as high-risk. Algorithmic decision-making didn't reduce bias: it systematised it, at scale, behind a veneer of objectivity. My argument: consistency is not fairness, and auditable does not mean fair."

Why it's better: - Names the system, the year, the context, the finding - States the thesis clearly and boldly - Grounds the abstract question in a real case from the course

---

Quick reference: argument checklist

Before recording, check:

- ☐ Am I arguing the **assigned side** (not my personal view)?
- ☐ Is my thesis **one clear sentence**?
- ☐ Does each supporting point have **specific evidence** (name, date, number)?
- ☐ Is my counterargument the **strongest** objection, not a weak one?
- ☐ Do I **acknowledge** what the counterargument gets right before responding?
- ☐ Am I speaking from **notes** (not reading a script)?
- ☐ Is my video **3–5 minutes**?

FIT1061 D2: Argumentation Guide